

lib_aec: Test Cases

Oscar Bailey

September 17, 2018

1 Method

To test the AEC, we must supply two audio inputs: a microphone input (AudioIn) and a far-end input (AudioRef). The microphone input is a combination of desired near-end audio and noise from the environment. After processing AudioIn and AudioRef, the AEC outputs two independent audio tracks: AudioOut, and a copy of the original AudioRef. AudioOut should be similar to the microphone input minus the noise produced by the far-end input. The noise produced by the far-end input will change depending on how the sound is reflected in the environment.

This document outlines test cases when the input for the AEC is simulated. The tests use a set of chosen parameters: Near-end Audio, AudioRef, and the transfer function of the room $H()$. AudioIn will be generated by summing the near-end audio with $H(\text{AudioRef})$. The AEC will be executed on AudioIn and AudioRef, producing $\text{AudioOut} = \text{AudioIn} - \hat{H}(\text{AudioRef})$ where $\hat{H}()$ is the model of $H()$ internal to the AEC. AudioOut will be compared to the near-end audio, producing a pass/fail result based on similarity within a chosen threshold.

Using simulated input allows us to focus only on cancelling noise created by AudioRef. When using real audio, $\text{AudioIn} = G(x) + H(y)$ where G is the transfer function describing the effect of the environment on near-end signal x . Therefore, a perfect AEC operating on recorded audio will generate $\text{AudioOut} = G(x)$. In comparison, by artificially constructing $\text{AudioIn} = x + H(y)$ a perfect AEC will generate $\text{AudioOut} = x$.

To pass a test, the AEC must produce an AudioOut that is sufficiently similar to the near-end audio. The method of comparison is illustrated in figure 3. Peaks in the frequency domain of the near-end audio should be present in AudioOut with a magnitude greater than -5dB relative to the near-end audio. Frequencies which are not present in the near-end audio should be less than -30dB in AudioOut relative to the largest magnitude in AudioOut.

For the first 60 seconds of processing, the AEC operates with silent near-end audio i.e. $\text{AudioIn} = H(y)$ and $\text{AudioRef} = y$. This ensures that the AEC has time to adapt its internal model of the room's transfer function and settle to a stable internal state.

An alternative test method is used to capture the impulse response of the transfer function internal to the AEC and compare it to the transfer function of the room. In this test, the input audio is silent followed by a single negative impulse. After the negative impulse, the AEC is processed without adaptation enabled. The resulting AudioOut is compared to the room transfer function. The test passes if they are similar.

When the transfer function of the room is longer than the length of the transfer function internal to the AEC, we expect that the AEC will not remove the far-end signal. In this case, the test fails if AudioOut is similar to the near-end audio.

We will test the performance of the AEC when using multiple channels on both the AudioIn and AudioRef inputs. In this case, each microphone has an independent model of the transfer function of the room. Each channel of the reference audio will have a different reverberation path from the speaker to the mic input. Therefore, for an M channel reference and an N channel microphone we will require $M \cdot N$ transfer functions to model the room.

As previously mentioned, the AEC outputs both AudioOut and a copy of AudioRef. We will compare the AudioRef output with the AudioRef input to test whether the AEC is distorting AudioRef while processing the input audio.

2 Construction

The transfer function of the room is visualised in figure 1. Table 1 describes test instances.

2.1 Reference Signal

These tests will use reference signals with varying spectral power distributions, and difference levels of auto-correlation. In preliminary tests, white noise as a reference signal has performed well. White noise has a wide,

uniform spectral power distribution and no autocorrelation.

We will also use reference signals which oscillate between a minimum and maximum frequency. The period of the oscillation will be greater than the excessive echo transfer functions, and therefore longer than the length of $\hat{H}()$. The two types of oscillation, discrete and continuous, are visualised in figure 2. The discrete 2-tone oscillating signal has some auto correlation and a small spectral power distribution. The continuous 2-tone signal has low auto correlation and a wide spectral power distribution.

To create a reference signal with a small spectral power distribution and high autocorrelation we will use a sine wave of constant frequency.

The reference signal output by the AEC will be compared sample-by-sample with the input reference signal. If they are not identical, the test will fail. This check will be performed for all tests.

2.2 Near-end Signal

The frequencies used for near-end audio will be within the range of human speech.

2.3 Headroom

To test for clipping and overflow errors during processing, the audio will be shifted right 2, 4, and 8 bits. The audio is stored as a sequence of 16 bit integers, so a right-shift by 8 bits removes half the precision. The spec dictates that at least 2 bits of headroom is required.

Test Name	HR Bits	Near-end Audio	Room	AudioRef
Simple Tests	{2,4,8}	Single frequency	{Small Echo, Long Echo, Decaying Echo}	{Discrete, Continuous, Single Sine, White Noise}
Multi-tone Tests	{2,4,8}	Multiple frequency	{Small Echo, Long Echo, Decaying Echo}	{Discrete, Continuous, Single Sine, White Noise}
Impulse Response Tests	{2,16}	Silence + Negative Impulse	{Small Echo, Long Echo, Decaying Echo}	{Discrete, Continuous, Single Sine, White Noise}
Excessive Echo Test	8	Single frequency	Excessive Echo	White Noise
Band Limited Test	4	Multiple frequency	Random	Band-limited White Noise

Table 1: Table of test instances.

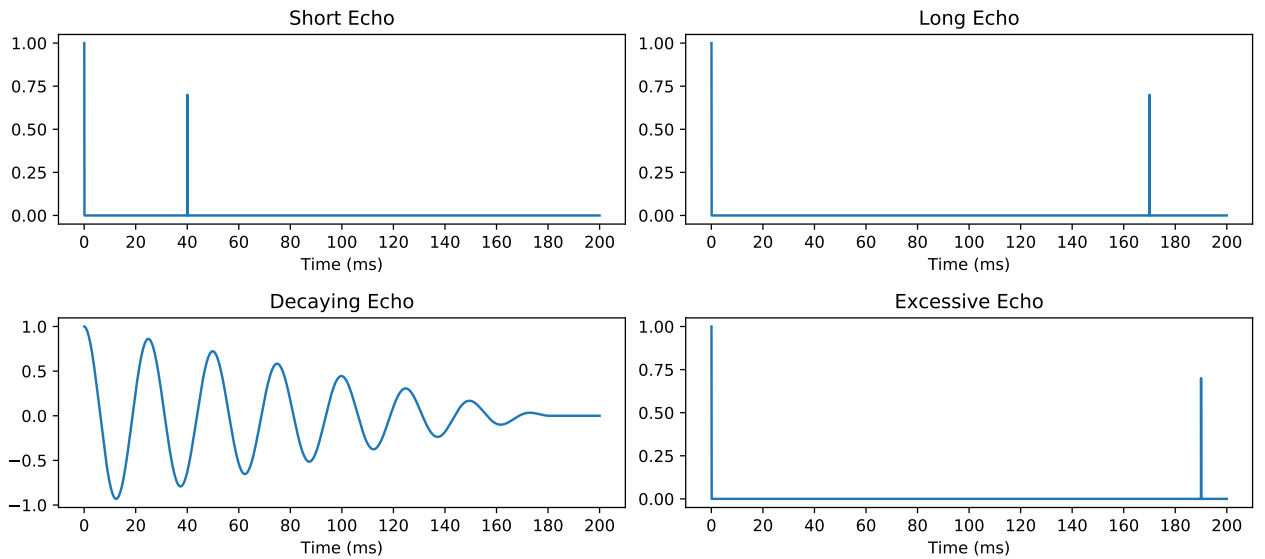


Figure 1: Impulse responses of different transfer functions used for testing.

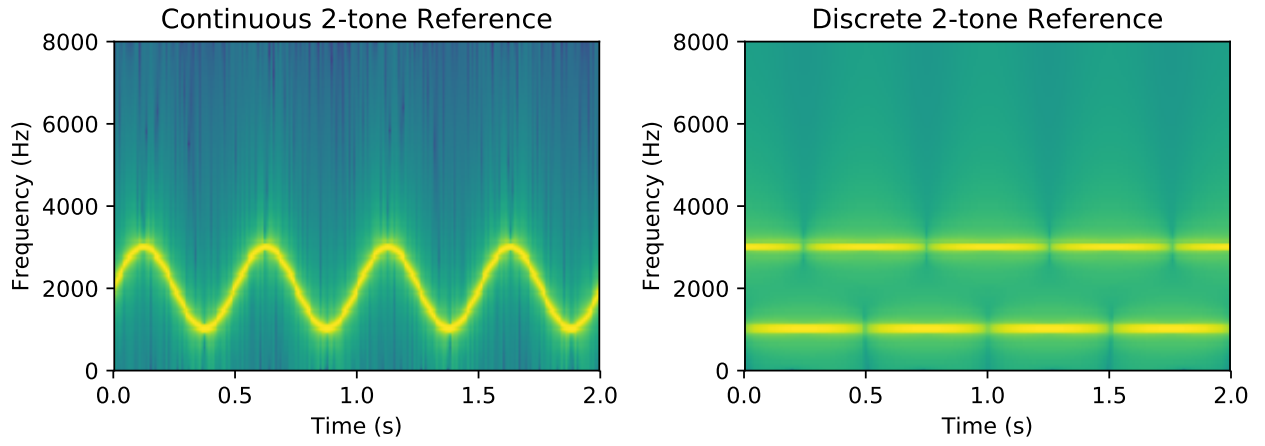


Figure 2: Discrete and Continuous versions of the 2-tone oscillating reference signal.




Figure 3: Frequency spectrum of single-frequency near-end audio. The red lines show where the frequency should be below 30dB, and the green lines show the frequency of greatest magnitude.